


Logical architecture

Compulsory Model 1 registry / GIS · Model 2 direct viewing · Model 3 VMS federation · Model 4 selected analytics




Statewide central recording is declined: $80,000 \times 2 \text{ Mbps} = 160 \text{ Gbps}$ (MODELLLED, docs/SCALE_MODEL.md).

District deployment is DESIGNED; VMS connectors are DEMO/TEST pending departmental access (docs/ADAPTERS.md).

A search, end to end


Four gates. Authentication, role, jurisdiction, and the purpose the search is bound to.



Purpose binding is not a dialog dismissed once — it is standing furniture in the interface, and it is written into every audit record the search produces (VERIFIED, tests/security/).

Evidence chain


Hash-chained, append-only, and honest about what it is not



A hash verifies unchanged bytes, not plate-to-image identity: older track-close stills can show a different vehicle.
Verify against source footage (docs/HLD.md §4.6). The certificate remains a draft, not a court determination.


Camera capability

Measured from the stream, per camera, per time band




Statewide target architecture

Video stays where it is recorded; metadata, events, requested clips and viewed streams move



DESIGNED (docs/STATEWIDE_ARCHITECTURE.md §2–§6; HLD §3, §6, §21). Registry and watchlist deltas, DR mirroring and evidence replication are in the boxes. Only inference compute grows in proportion to cameras analysed (MODELLED, reports/capacity_model.json). Nothing here ran at 80,000 cameras.

Detected and matched in the cell; the plate lane builds the route at the state; everything else goes to the lake



Lanes, lake, plate index and SLOs are DESIGNED (docs/STATEWIDE_ARCHITECTURE.md §5, §6, §9, §18b).